



IBM

# Linux Kernel: An Introduction

---

Ankita Garg ([ankita@in.ibm.com](mailto:ankita@in.ibm.com))  
RealTime Linux Team  
IBM Linux Technology Center

16/02/2008

Linux is a registered trademark of Linus Torvalds.

IBM



# IBM

## Linux Technology Center

- World-wide distributed team working with open-source communities
- Mission
  - Make Linux better
  - Accelerate its growth as an enterprise OS
  - Expand Linux Reach
  - Make it mature OS ready for mission critical workloads
- LTC India
  - Development
  - Internal support
  - Testing

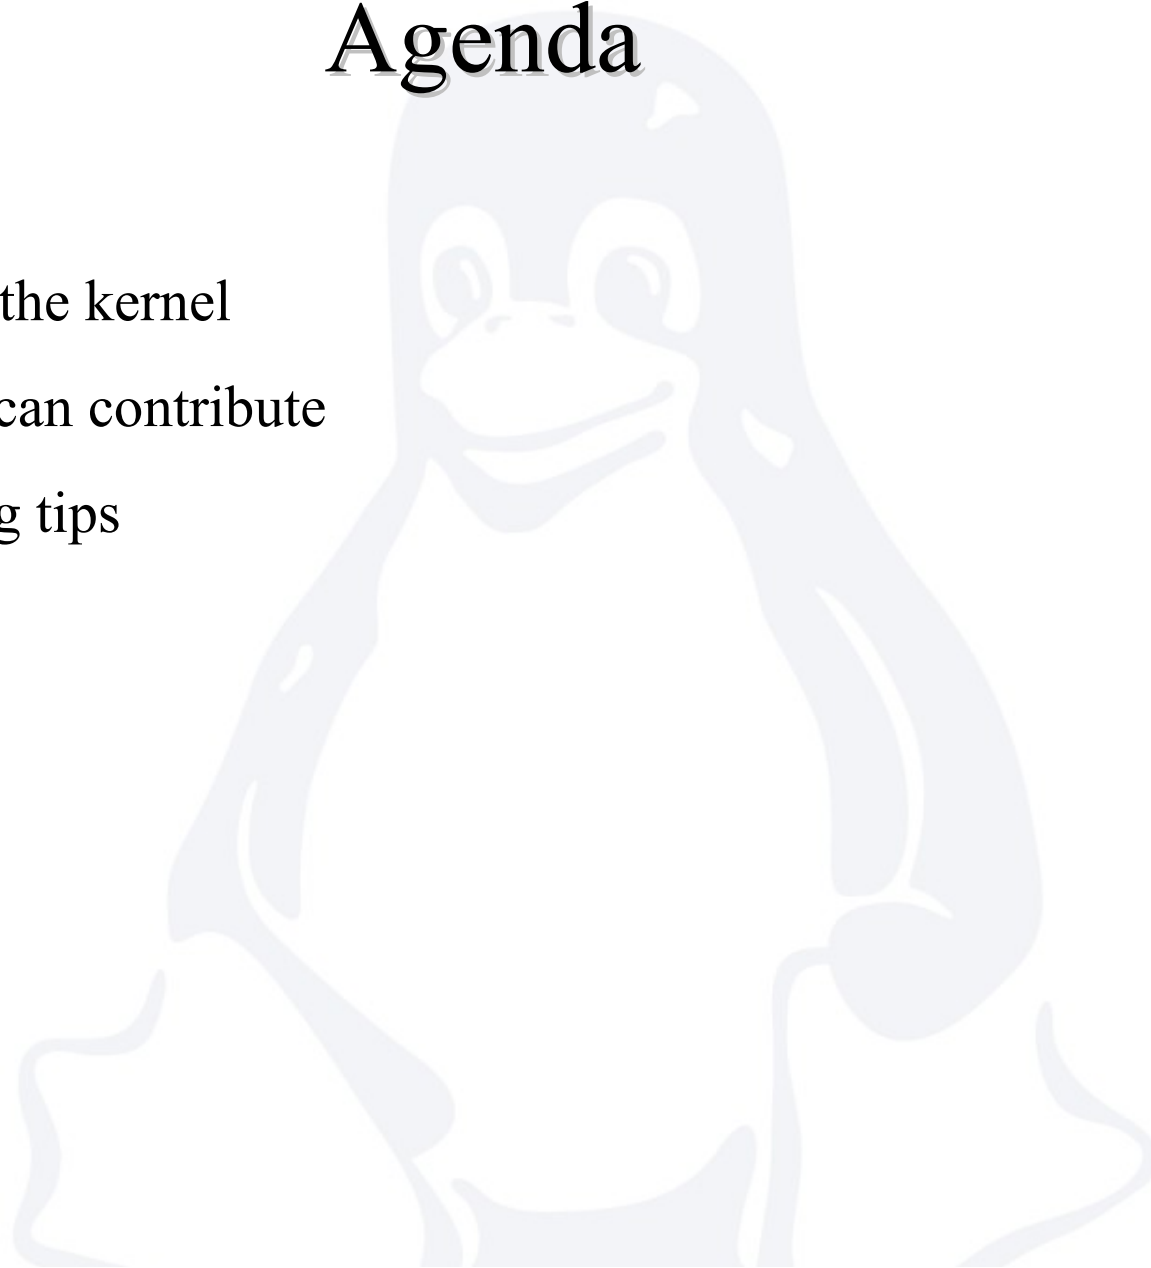
# Objective

Introduce the Linux kernel and the development community.


Inspire you to play with the kernel and contribute towards making it better

# Agenda

- Play with the kernel
- How you can contribute
- Debugging tips

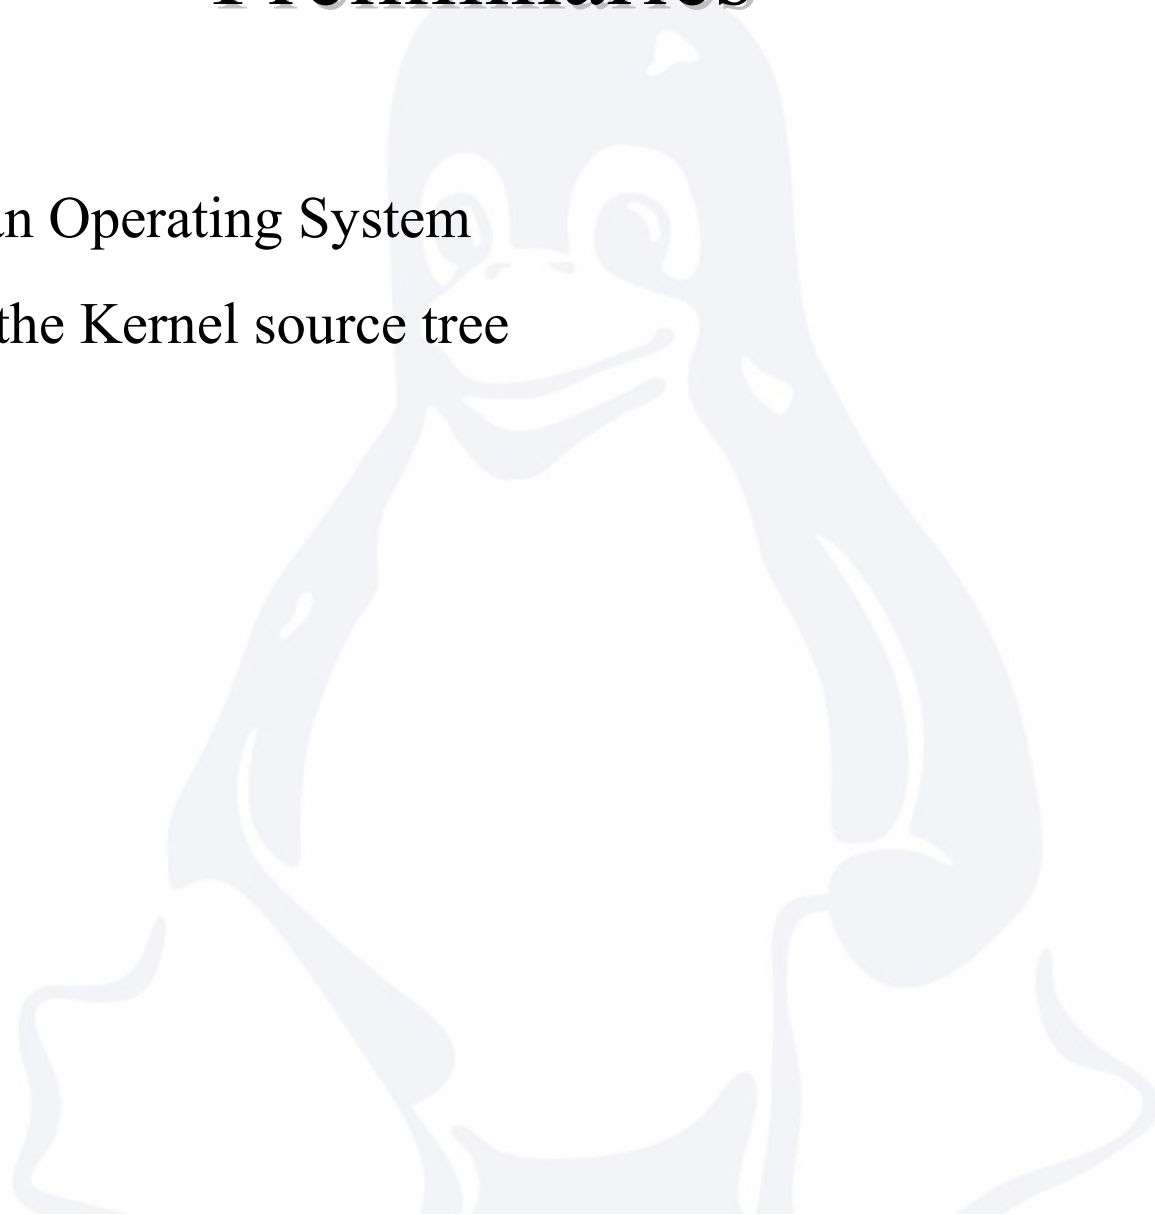


# Preliminaries



# Preliminaries

- Linux as an Operating System
- A look at the Kernel source tree

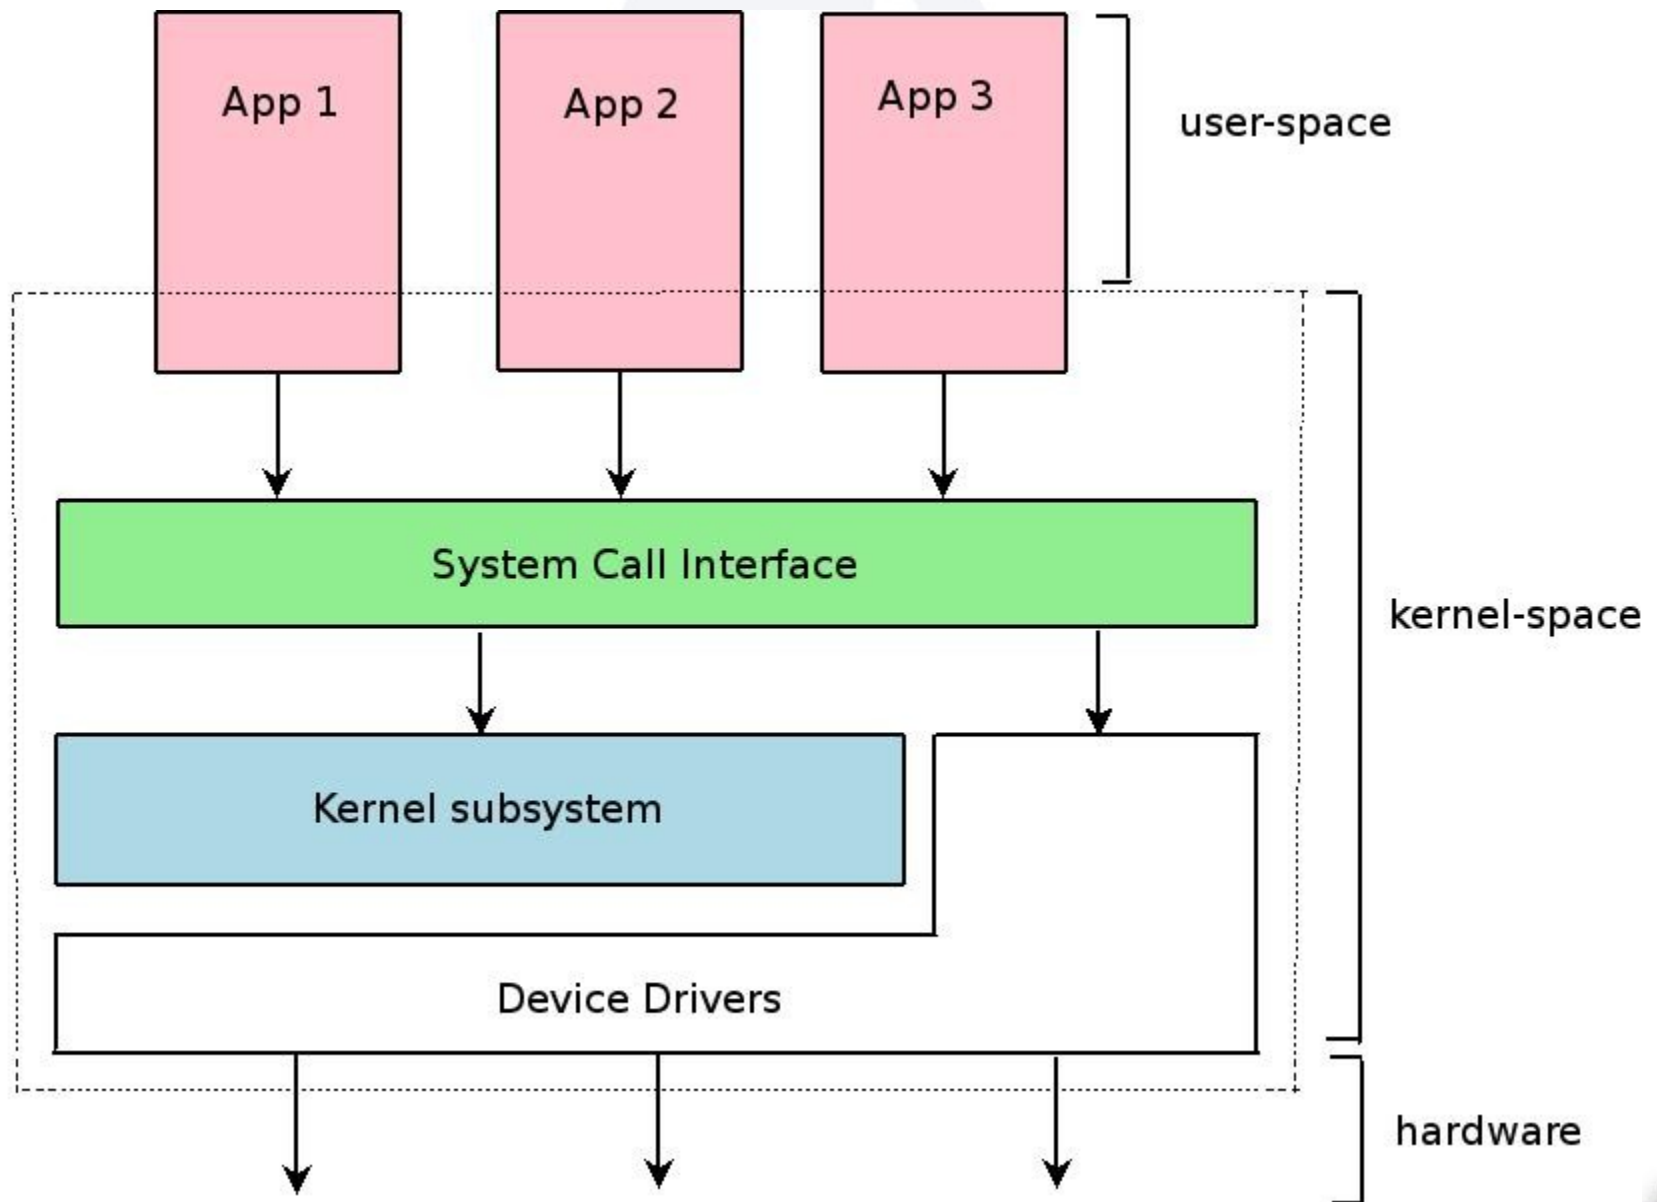




# What is an Operating System??




# Application, the kernel and hardware





# Channels

- Hardware dependency
- A beast of a different kind
- Linux distributions — RedHat, SUSE, Fedora, Debian, Ubuntu, Mandriva, YellowDog Linux, Puppy Linux, Gentoo, Slackware Linux

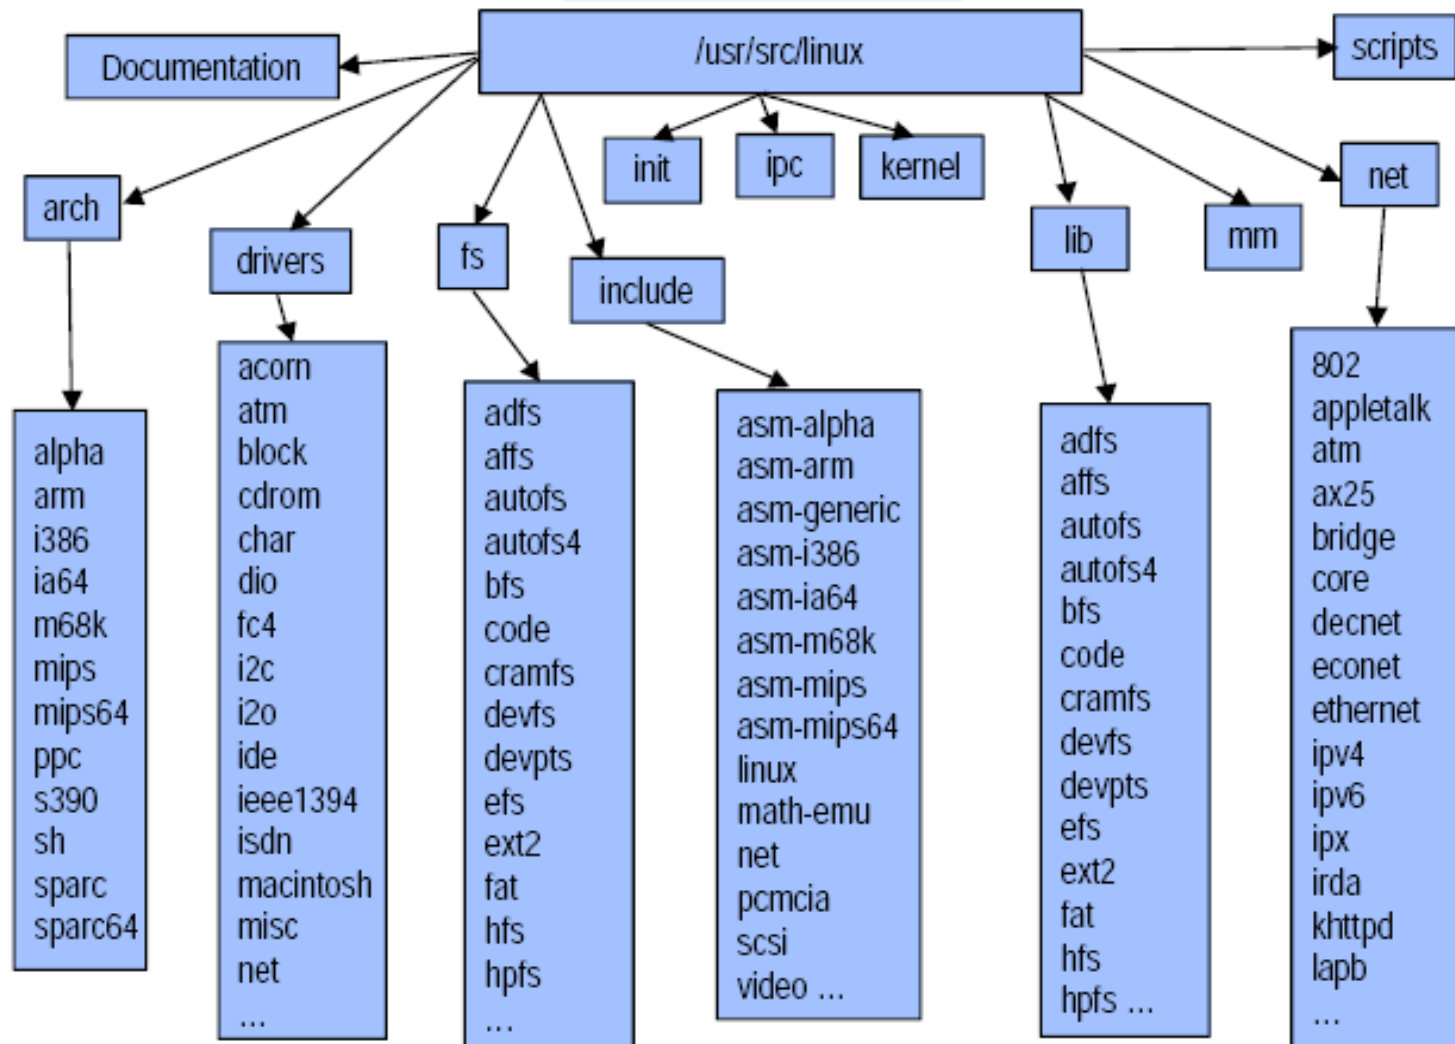


# Obtaining the Kernel Source

- ftp from kernel.org
  - `ftp ftp.<country>.kernel.org`
- clone the linux git tree
  - `git clone git://address/linux-2.6.git linux-2.6`
- Install kernel source from the distribution CDs
  - `rpm -ivh <path-to-src-rpm>/kernel-src.rpm`



# Getting Started – Kernel Source layout





Let's build our own kernel

# Configuration

- First step is to configure the kernel

```
$ make config  
$ make oldconfig  
$ make menuconfig  
$ make gconfig  
$ make xconfig  
$ make defconfig
```

- Resultant .config file
- Tricky
- Tips
  - /proc/cpuinfo for type of cpu
  - lspci -v
  - dmesg | less
  - enable module loading

# Compilation

- Build core kernel image

*\$ make -j<no of jobs> bzImage*

- Build and install modules

*\$ make modules*

*\$ make modules\_install*

- Recompile kernel

*\$ make clean*

*\$ make mrproper*



# Install & boot into new kernel

- Copy kernel to required location

```
$ cp arch/i386/boot/bzImage /boot/vmlinuz-2.6.24
```

```
$ cp System.map /boot/System.map-2.6.24
```

- Edit the bootloader config file

- Grub: Edit /etc/grub.conf or /boot/grub/menu.lst
- LILO: /etc/lilo.conf
- Yaboot: /etc/yaboot.conf

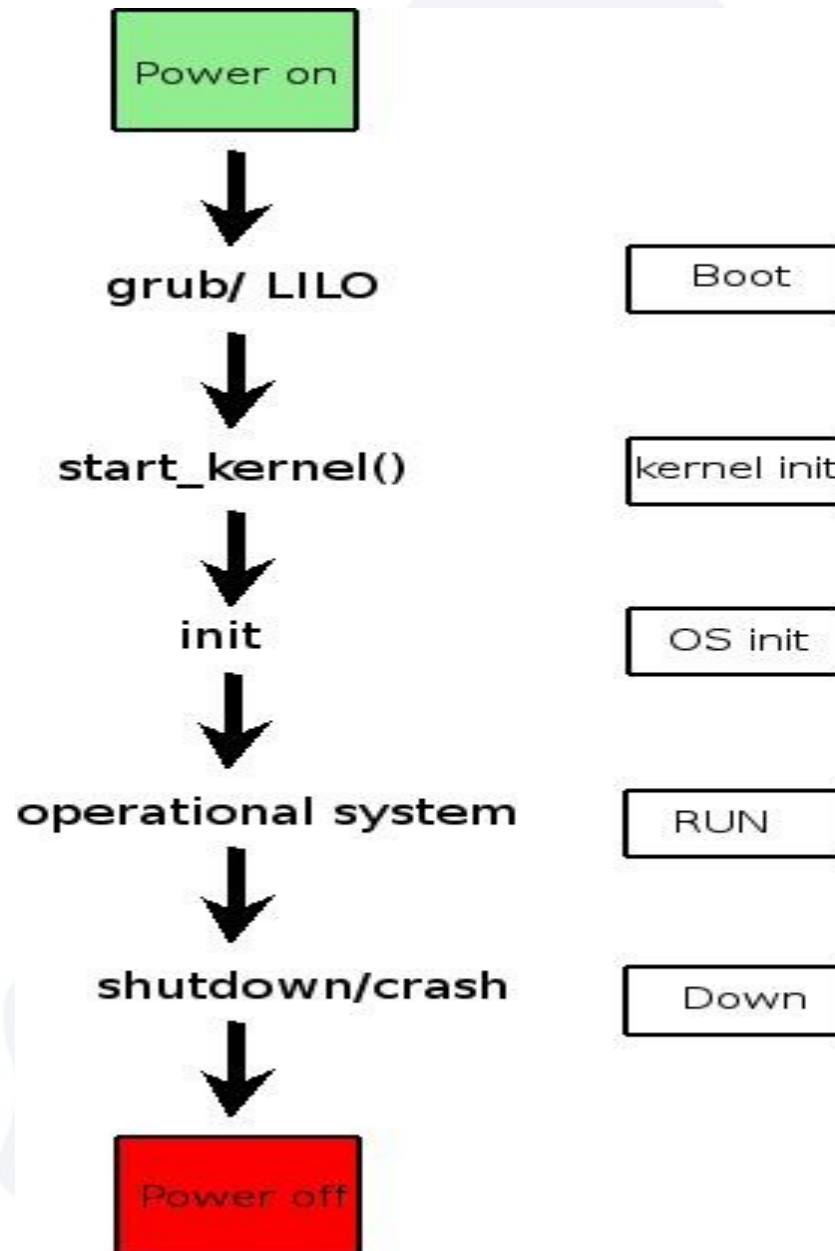
- Recompile kernel

```
$ make clean
```

```
$ make mrproper
```

- Install kernel – manually or using a command

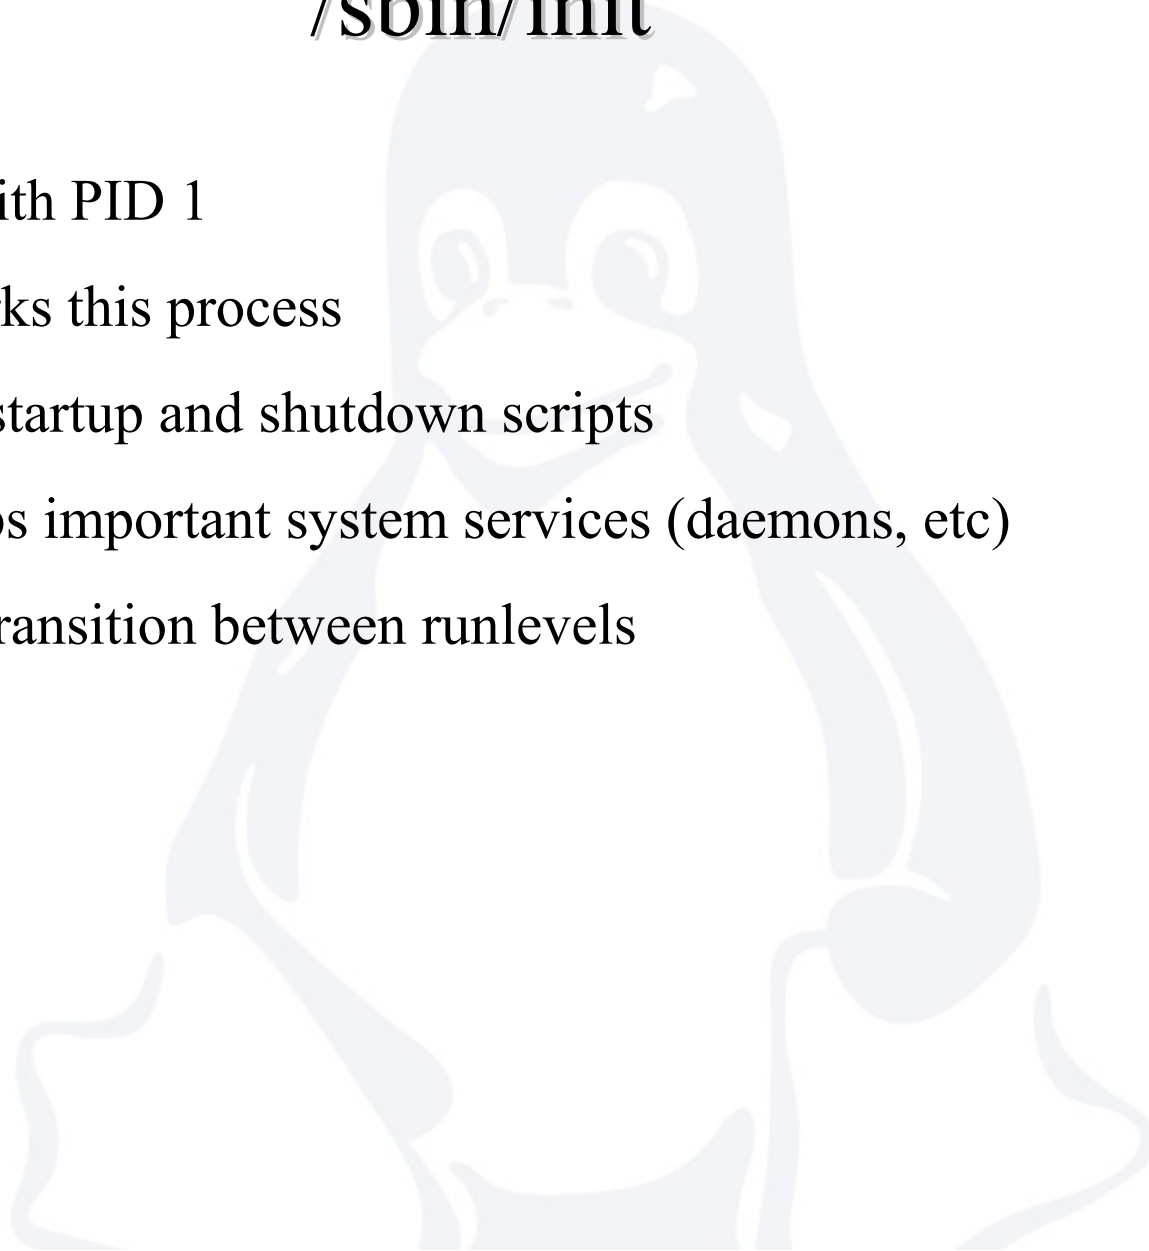
# Boot up and kernel initialization





# /sbin/init

- Process with PID 1
- Kernel forks this process
- Executes startup and shutdown scripts
- Starts/stops important system services (daemons, etc)
- Controls transition between runlevels





Using printk





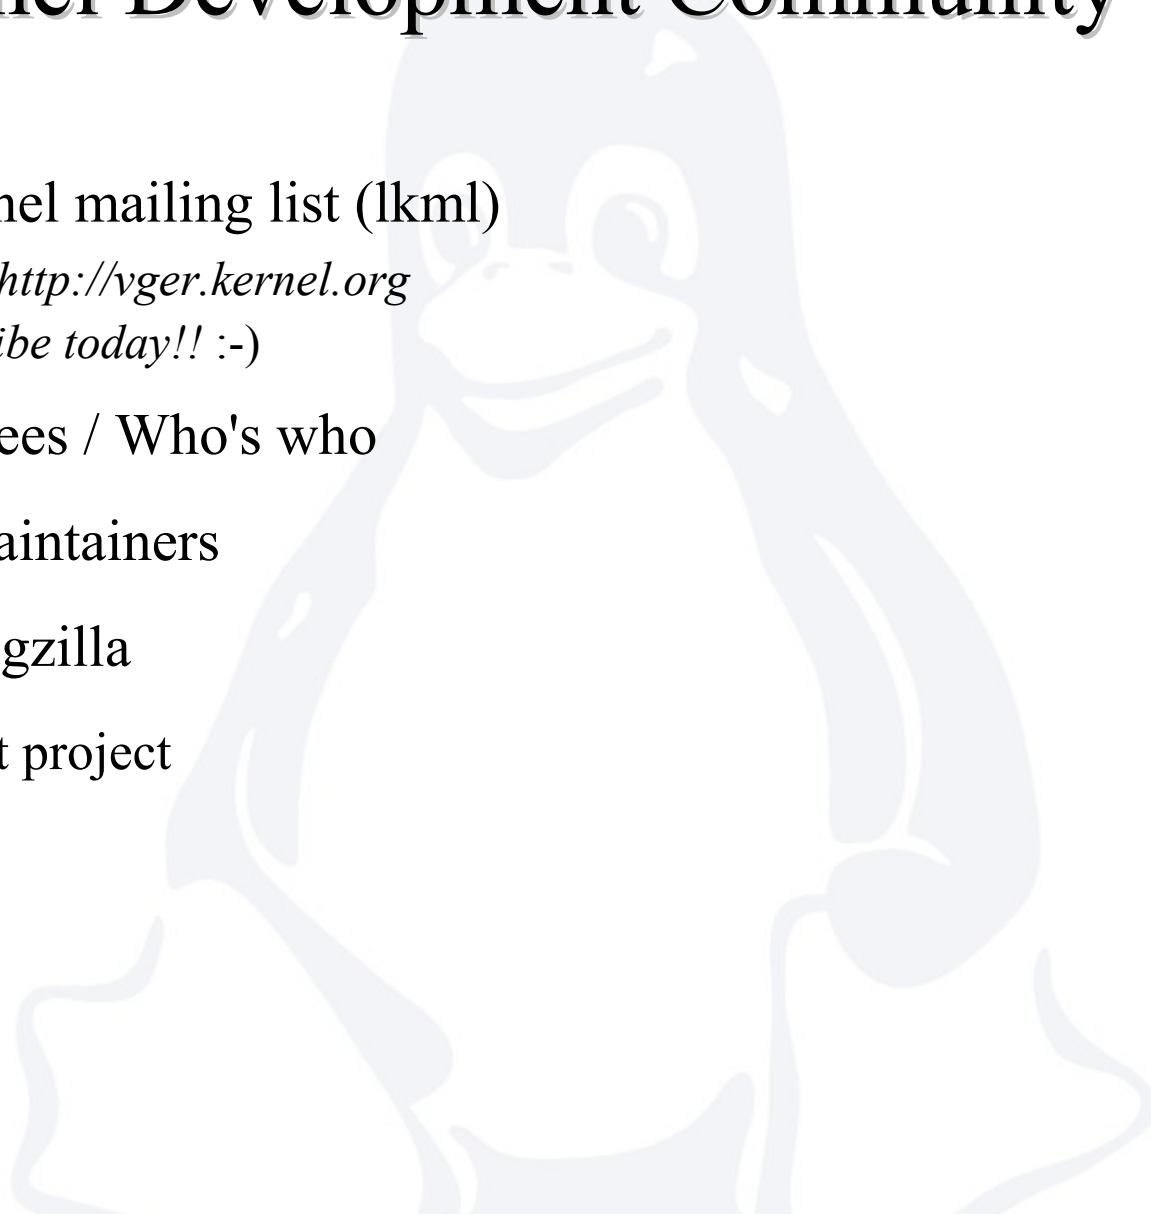
# How you can contribute

IBM



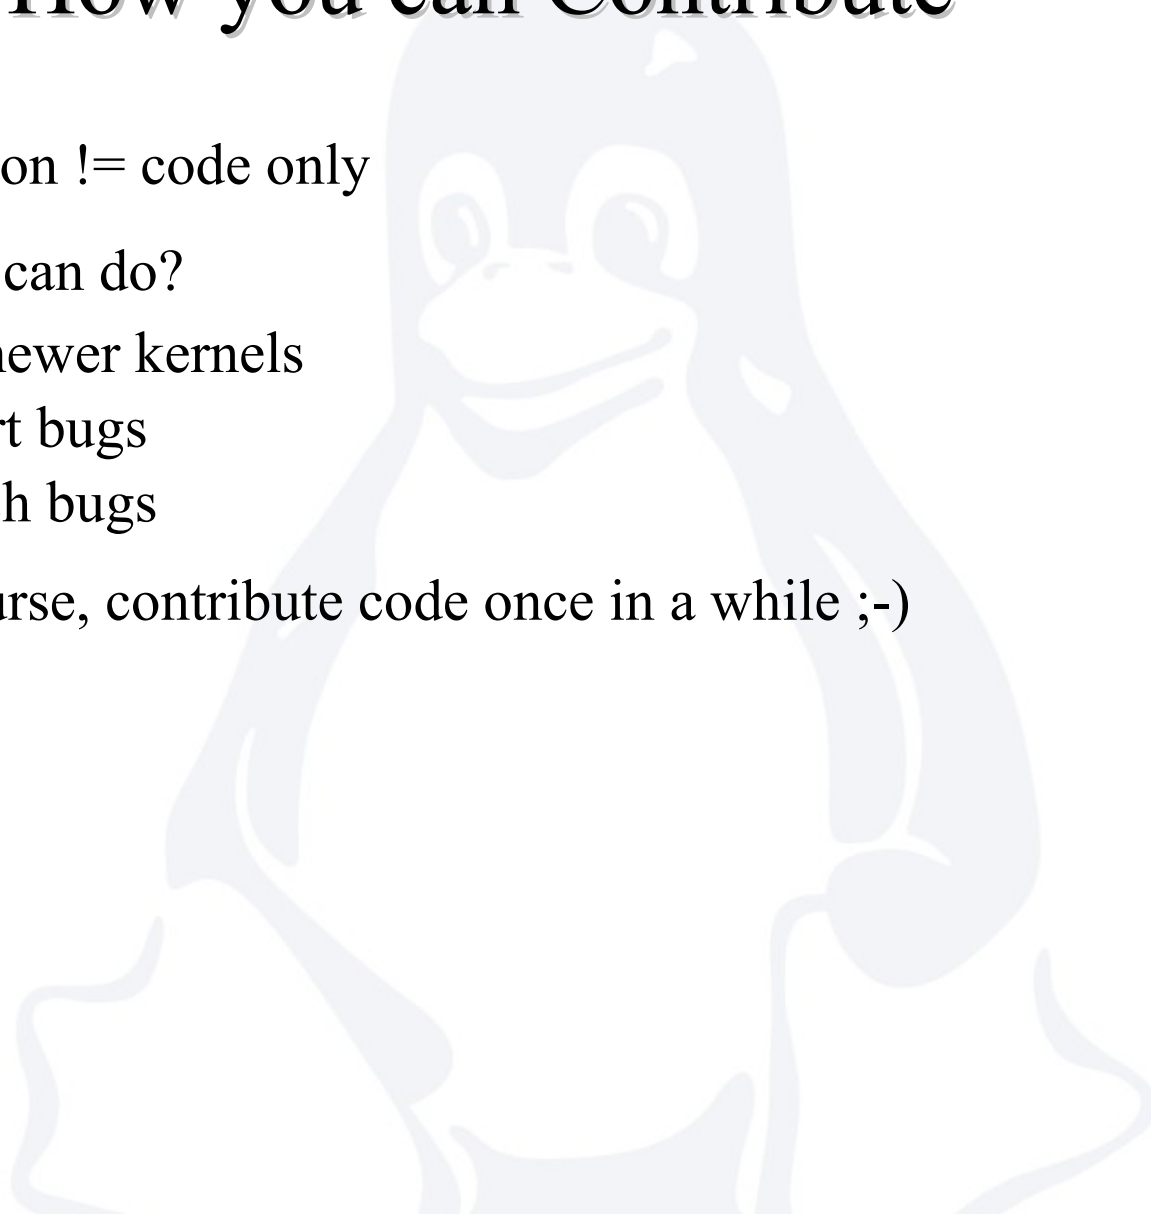
# Kernel Development Community

- Linux kernel mailing list (lkml)
  - *Info@ <http://vger.kernel.org>*
  - *Subscribe today!! :-)*
- Various trees / Who's who
- Role of maintainers
- Kernel Bugzilla
- Kernel Test project



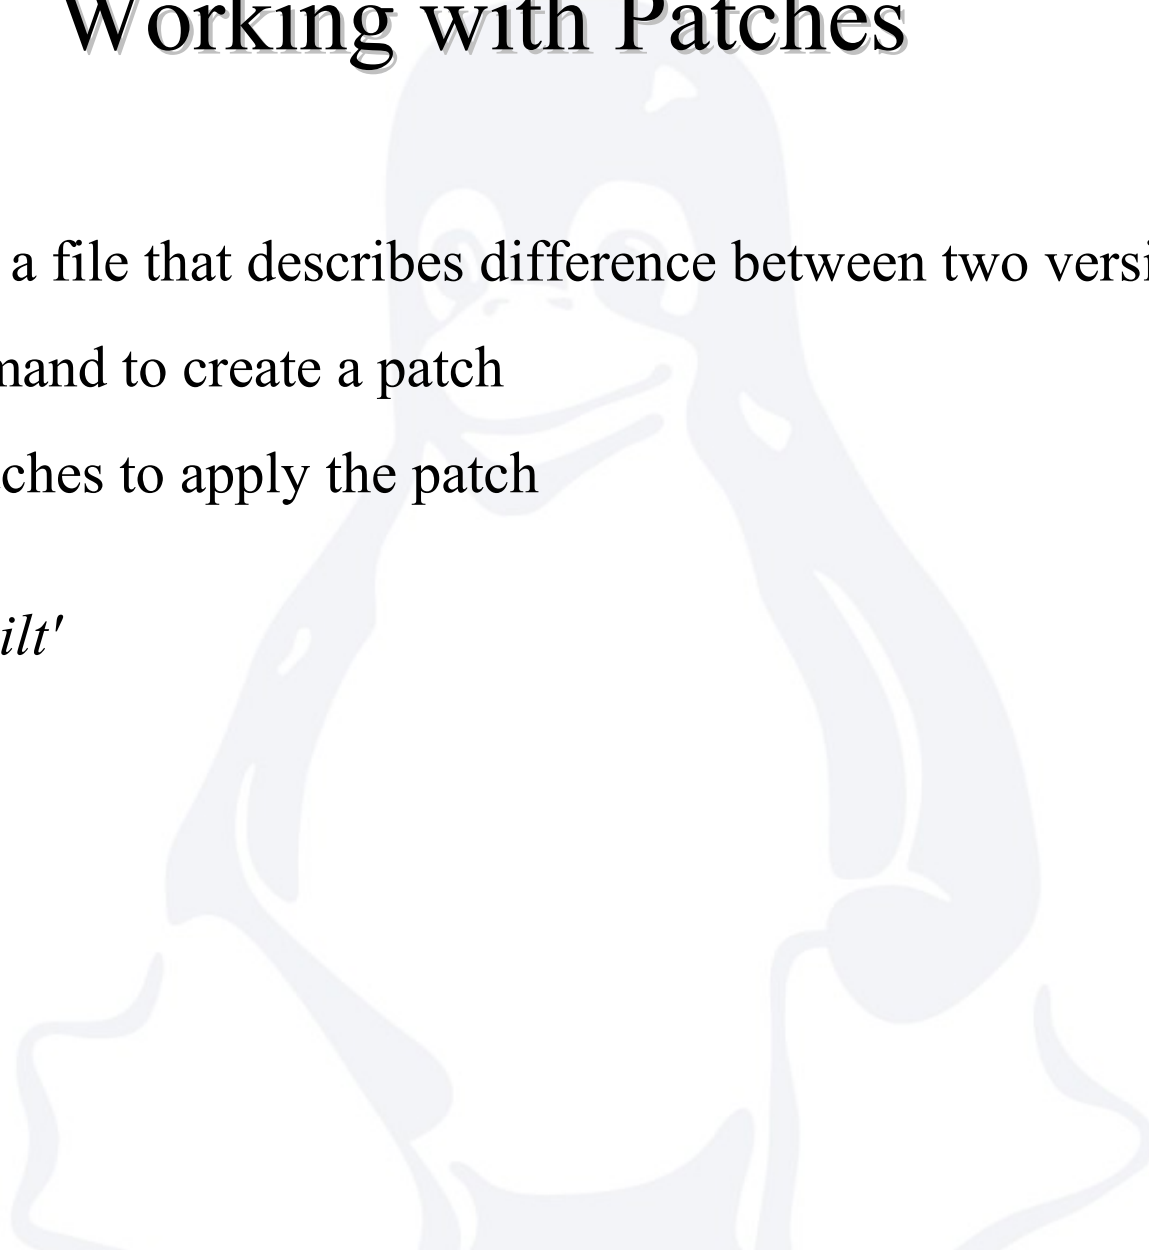
# How you can Contribute

- Contribution != code only
- What you can do?
  - Test newer kernels
  - Report bugs
  - Squash bugs
- And ofcourse, contribute code once in a while ;-)



# Working with Patches

- A patch is a file that describes difference between two versions of a file
- *'diff'* command to create a patch
- *'patch'* patches to apply the patch
- Or use *'quilt'*



# Code contribution

- Refer kernel Documentation (/Documentation)
- Create Patches
- Test the patch, get reviews
- Dos & Dont's
  - Read through the patch, check if unwanted files included in the patch
  - watch the level of directories
  - create against latest kernel version
  - Coding Style
- Submitting Patches
  - Refer the MAINTAINERS file (though most often out-of-date)
  - Watch the mailing list
  - Ask somebody
- Distributing your patches



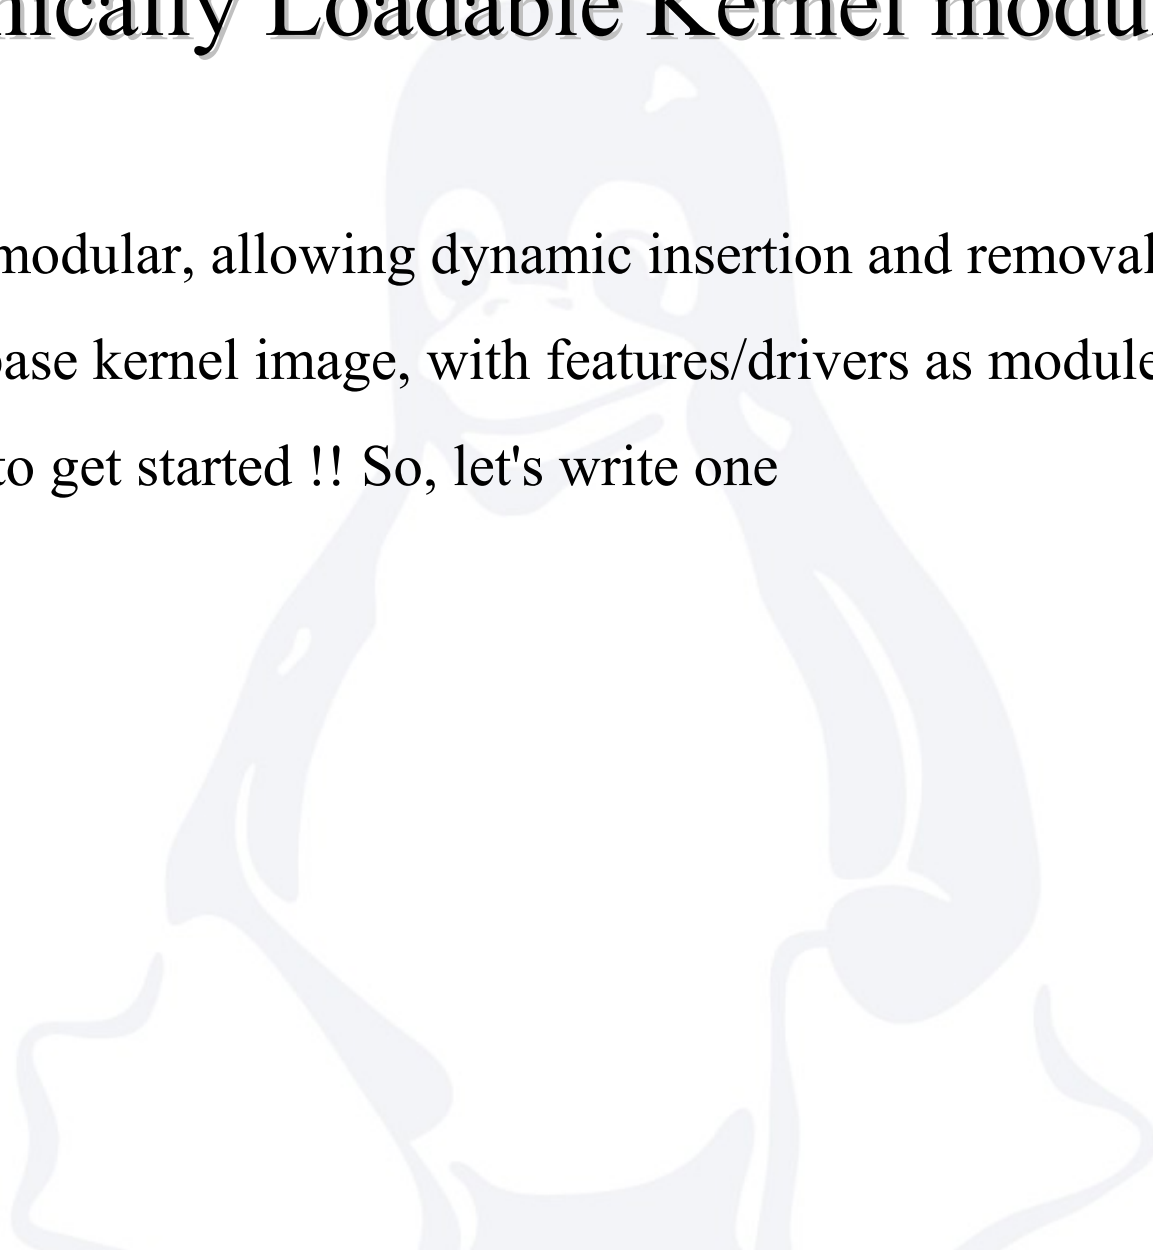
# Writing kernel modules





# Dynamically Loadable Kernel module

- Kernel is modular, allowing dynamic insertion and removal of code
- Minimal base kernel image, with features/drivers as modules
- Best way to get started !! So, let's write one



# myModule.c

```
/* myModule.c */

#include<linux/module.h>
#include<linux/init.h>
#include<linux/kernel.h>

static int __init myhello_init(void)
{
    printk(KERN_INFO "\n Hello there, mymodule loaded \n");
    return 0;
}

static void __exit myhello_exit(void)
{
    printk(KERN_INFO "\n Have a nice day! myModule unloaded \n");
}

module_init(myhello_init);
module_exit(myhello_exit);

MODULE_LICENSE("GPL");
MODULE_AUTHOR("XYZ");
```



# Building & installing myModule.c

- Edit Makefile in the same directory..add the following line:

*obj-m += myModule.o*

- Compile module

*# make -C <top level kernel src tree> SUBDIRS=\$PWD modules*

- Load the module

*# insmod myModule.ko*

- Verify module loaded

*# dmesg | tail*



# Building & installing a module

- Write a makefile

```
obj-m := kprobe-example.o
```

```
KDIR := /lib/modules/$(shell uname -r)/build
```

```
PWD := $(shell pwd)
```

```
default:
```

```
$(MAKE) -C $(KDIR) SUBDIRS=$(PWD) modules
```

```
clean:
```

```
rm -f *.mod.c *.ko *.o
```

- Build using *make*
- Install using *insmod*

# Linux Modules

- To make it part of the kernel source
- Module Parameters

```
module_param(name, type, perm);
```

- Pass as *name=value* while inserting
- Use lsmod, insmod, modprobe, rmmod



# Debugging Methods/Tips



# Debugging Tips

- Use printk
- /proc filesystem
- kprobes / SystemTap (link to stap)
- KDB
- KGDB
- QEMU
- kdump
- lockstat, logdev



# Useful Tools

- Source code browsing – cscope, lxr, ctags, etc
  - Cscope
    - <http://cscope.sourceforge.net/>
    - [http://cscope.sourceforge.net/cscope\\_vim\\_tutorial.html](http://cscope.sourceforge.net/cscope_vim_tutorial.html)
  - LXR
    - <http://lxr.linux.no/>
- Patching the kernel – patch, quilt
- GIT





# Start Contributing

- Read/Understand the kernel source
- [www.kernelnewbies.com](http://www.kernelnewbies.com)
- [lwn.net](http://lwn.net)
- Linux kernel Documentation
- Books
  - *Linux Kernel Development* by Robert Love
  - *Understanding the Linux Kernel* by Bovet and Cesati
  - *Linux Device Drivers* by Alessandro Rubini
- Follow the mailing list
- Use Linux as your primary OS



A large, faint, light blue illustration of a penguin, likely Tux, is centered in the background. The penguin is standing and facing forward, with its arms slightly out to its sides.

# Questions ??



\*\*\*\*\*

## Legal Statement

This work represents the views of the author(s) and does not necessarily reflect the views of IBM Corporation.

Linux is a registered trademark of Linus Torvalds.

Other company, product, and service names may be trademarks or service marks of others.

\*\*\*\*\*



Thank You !!

IBM



# Back up Slides





A tour of [www.kernel.org](http://www.kernel.org)

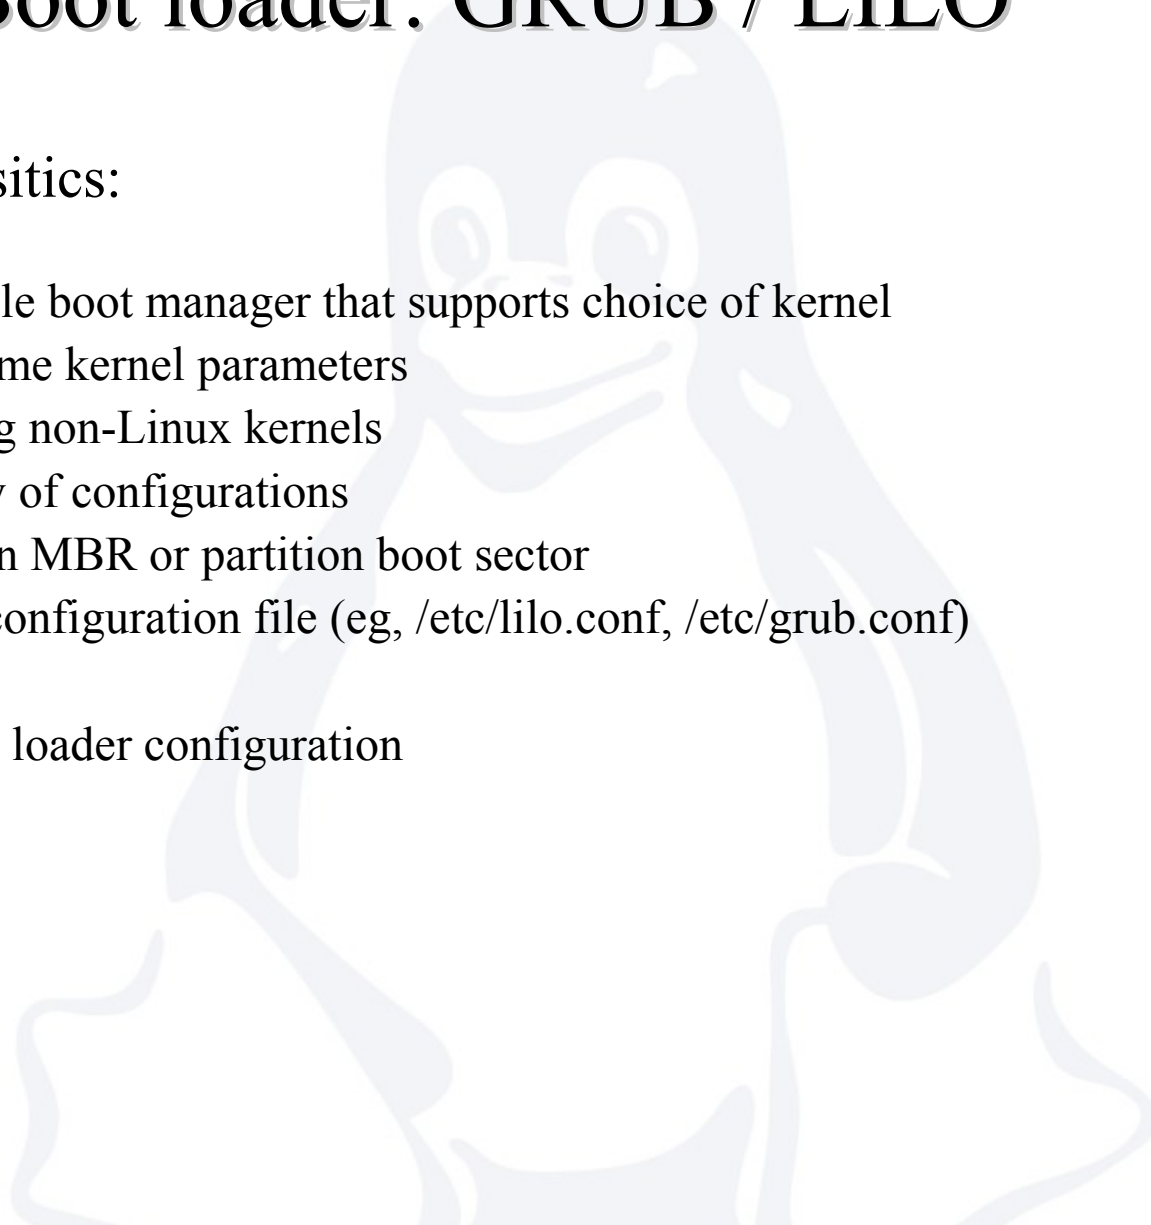


# Boot loader: GRUB / LILO

## ■ Charactersitics:

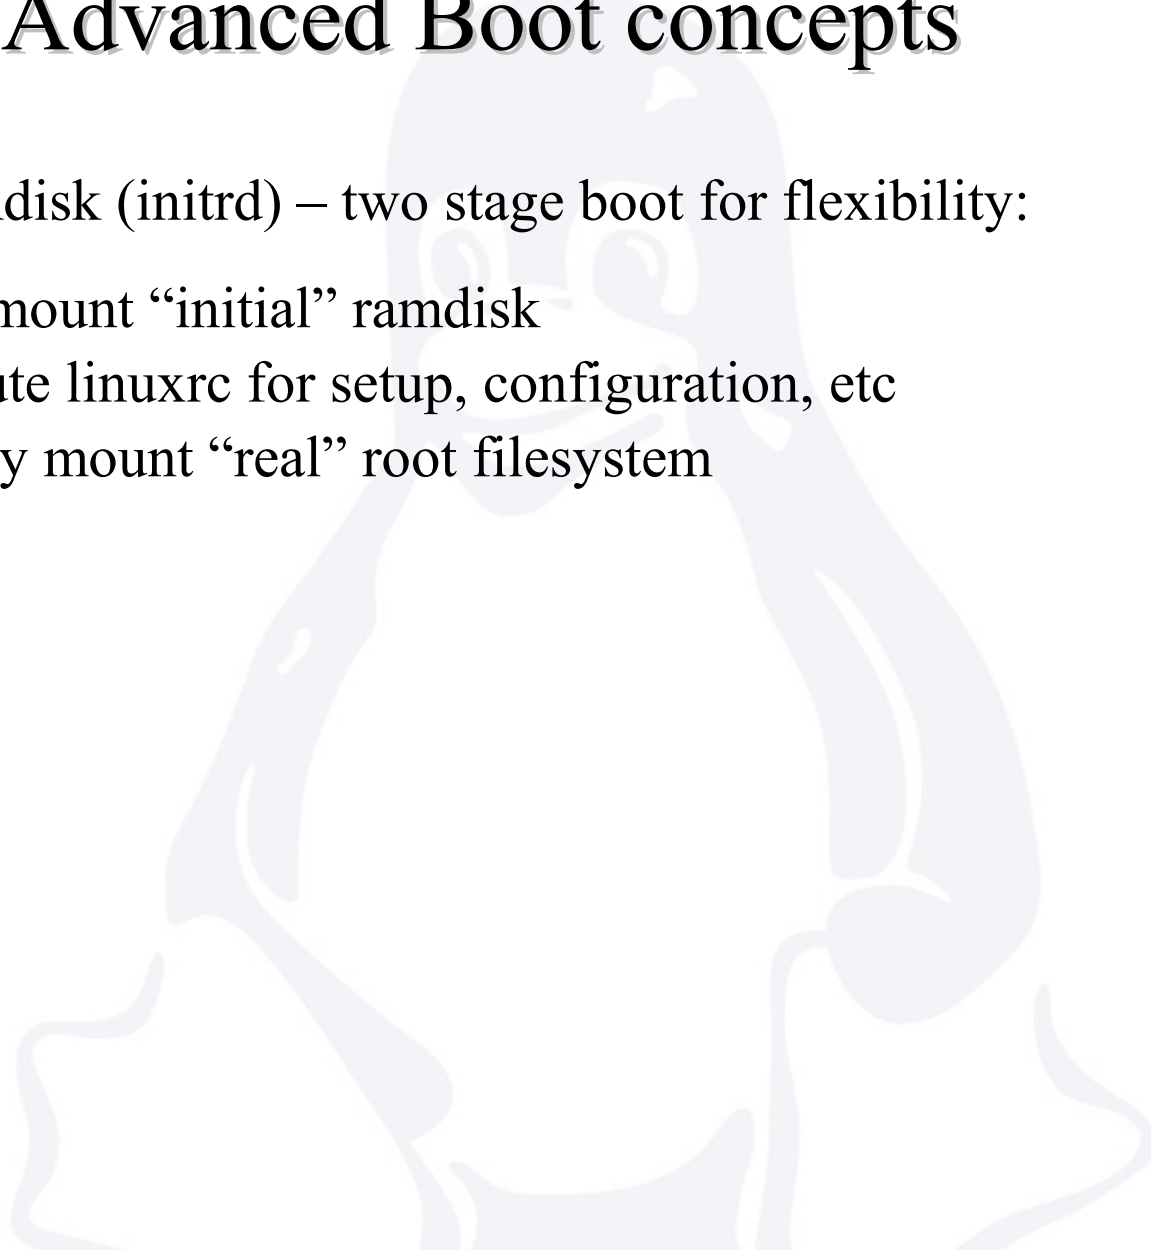
- Versatile boot manager that supports choice of kernel
- Boot time kernel parameters
- Booting non-Linux kernels
- Variety of configurations
- Lives in MBR or partition boot sector
- Has a configuration file (eg, /etc/lilo.conf, /etc/grub.conf)

## ■ Sample boot loader configuration



# Advanced Boot concepts

- Initial ramdisk (initrd) – two stage boot for flexibility:
  - First mount “initial” ramdisk
  - Execute linuxrc for setup, configuration, etc
  - Finally mount “real” root filesystem





# /proc filesystem

- virtual filesystem mounted under /proc
- Provided information on running processes
  - read-only access to kernel data structures
  - superuser can change kernel parameters at runtime
- [Linux kernel Documentation](#)

